



SCHOOL OF COMPUTATION,  
INFORMATION AND TECHNOLOGY —  
INFORMATICS

TECHNISCHE UNIVERSITÄT MÜNCHEN

Master's Thesis, ... in Informatics

**Hiding Cache-Misses with Coroutines  
across different Hardware Architectures**

Daniel Gruber





SCHOOL OF COMPUTATION,  
INFORMATION AND TECHNOLOGY —  
INFORMATICS

TECHNISCHE UNIVERSITÄT MÜNCHEN

Master's Thesis, ... in Informatics

**Hiding Cache-Misses with Coroutines  
across different Hardware Architectures**

**Verbergen von Cache-Misses mittels  
Coroutinen auf unterschiedlichen  
Hardwarearchitekturen**

Author:	Daniel Gruber
Examiner:	Alexander Beischl
Supervisor:	Prof. Dr. Thomas Neumann
Submission Date:	22.12.2025



I confirm that this master's thesis, ... is my own work and I have documented all sources and material used.

Munich, 22.12.2025

Daniel Gruber

## **Acknowledgments**

# **Abstract**

# Contents

<b>Acknowledgments</b>	<b>iv</b>
<b>Abstract</b>	<b>v</b>
<b>1 Introduction</b>	<b>1</b>
1.1 Section . . . . .	1
<b>2 Coroutines, Data Structures and Cache Misses Fundamentals</b>	<b>2</b>
2.1 C++ Coroutines . . . . .	2
2.1.1 The co_await, co_return and co_yield operators . . . . .	4
2.1.2 Example of Coroutines . . . . .	4
2.1.3 Simple generator (co_yield) example . . . . .	4
2.1.4 Tricks and Pitfalls . . . . .	4
2.2 Computer Architecture, Cache Misses . . . . .	4
2.2.1 Different Computer Architectures, x86-64, amd, apple, mips . . . . .	4
2.2.2 Introduction to computer achitecture and storage layout . . . . .	4
2.2.3 What are cache misses? . . . . .	4
2.2.4 When can they happen? . . . . .	4
2.2.5 Why are they relevant in data bases and different index structures	4
2.3 Data Structures . . . . .	4
2.3.1 Hashtable . . . . .	4
2.3.2 B+ Tree . . . . .	4
<b>3 Using Coroutines in different DataStructures for hiding cache-misses</b>	<b>5</b>
3.1 Different coroutine approaches and their trade-offs . . . . .	5
3.2 Coroutines in chaining Hashtable . . . . .	5
3.2.1 Implementation . . . . .	5
3.2.2 Benchmarking and Measurements . . . . .	5
3.2.3 Results . . . . .	5
3.3 Coroutines in linear probing hashtable . . . . .	5
3.3.1 Implementation . . . . .	5
3.3.2 Measurements . . . . .	5
3.3.3 Results . . . . .	5

*Contents*

---

3.4 Coroutines in B+ Tree . . . . .	5
3.4.1 Implementation . . . . .	5
3.4.2 Measurements . . . . .	5
3.4.3 Results . . . . .	5
<b>4 Integration of Coroutines in DuckDB</b>	<b>6</b>
4.1 DuckDB Introduction . . . . .	6
4.2 Understanding the implemented linear Hashtable . . . . .	6
4.3 Integration of Coroutines in DuckDB . . . . .	6
4.4 Results of timings . . . . .	6
<b>5 Evaluation</b>	<b>7</b>
5.1 Coroutine Approaches evaluation . . . . .	7
5.2 Coroutine in different Data structures evaluation . . . . .	7
5.3 Coroutine in DuckDB evaluation . . . . .	7
5.4 Summarized evaluation of coroutines in C++ for hiding cache misses .	7
<b>6 Conclusion</b>	<b>8</b>
<b>Abbreviations</b>	<b>9</b>
<b>List of Figures</b>	<b>10</b>
<b>List of Tables</b>	<b>11</b>
<b>Bibliography</b>	<b>12</b>

# 1 Introduction

## 1.1 Section

Citation test [4].

How can Coroutines hide latencies for hashtable/bptree lookups? What is the performance gain generally? What is the performance gain in duckdb?

How does the performance gain differ between various architectures? (Intel x86 vs ARM vs NUMA) What about Threading and Coroutines? Hyperthreading/NUMA

## 2 Coroutines, Data Structures and Cache Misses Fundamentals

### 2.1 C++ Coroutines

C++ Coroutines are available from the C++20 standard on, and from the C++23 standard there is a generator based on coroutines.

A coroutine is a generalization of a function, being able additionally to normal functions to suspended execution and resumed it later on. Any function is a coroutine if its definition contains at least one of these three keywords:

- `co_await` - to suspend execution
- `co_yield` - to suspend execution and returning a value
- `co_return` - to complete execution and returning a value

C++ Coroutines are stackless, meaning by suspension and consecutively returning to the caller, the data of the coroutine is stored separately from the stack, namely on the heap.

To illustrate the difference between coroutines and functions, this paragraph provides a background and general information of function calls and return: A normal function has a single entry point - the Call operation - and a single exit point - the Return operation. The Call operation creates an activation frame, suspends execution of the caller and transfers execution to the callee, where the caller is the invoking function and the callee is the invoked function. The Return operation returns the value in the return statement to the caller, destroys the activation frame and then resumes execution of the caller. These operations include calling conventions splitting the responsibilities of the caller and callee regarding saving register values to their activation frames. The activation frame is also commonly called stack frame, as the functions state (parameters, local variables) are stored on the stack. Normal Functions have strictly nested lifetimes, meaning they run synchronously from start to finish, allowing the stack to be a highly efficient memory allocation data-structure for allocation and freeing frames. The pointer pointing at the top of the stack is the `**rsp**` register on X86-64 CPU Architectures.

Coroutines have, additionally to the call and return operation, three extra operations, namely suspend, resume and destroy. As coroutines can suspend execution without destroying the activation frame, as it may be resumed later, the activation frames are not strictly nested anymore. This requires that the coroutines state is saved to the heap, like illustrated in Figure 2.1 where a function f() calls a coroutine c().

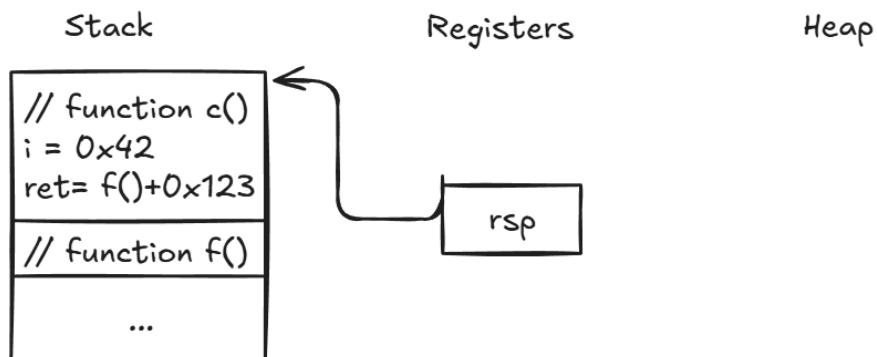


Figure 2.1: Calling a Coroutine

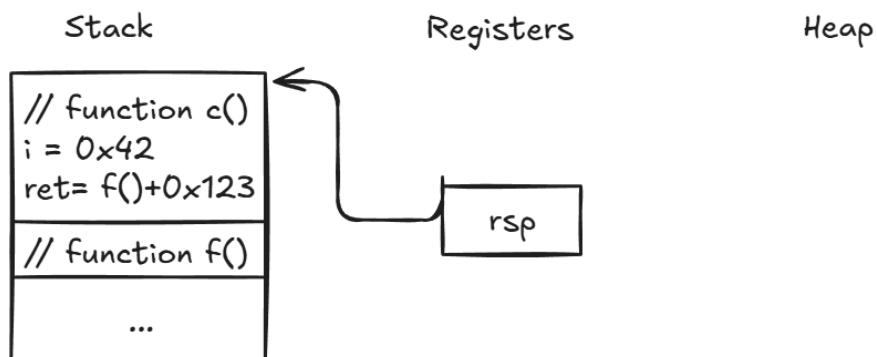


Figure 2.2: Calling a Coroutine

**AWAITABLE** An Awaitable has to provide the following three methods:

```
bool await_ready();

// one of:
void await_suspend(std::coroutine_handle<> caller_of_co_await);
bool await_suspend(std::coroutine_handle<> caller_of_co_await);
std::coroutine_handle<> await_suspend(std::coroutine_handle<> caller_of_co_await);
```

```
T await_resume();
```

Listing 2.1: awaitable type

PROMISE TYPE

[1] [3] [2]

### 2.1.1 The `co_await`, `co_return` and `co_yield` operators

`co_await`

awaitable (and awaiter) concept and promiseType

`co_return`

`co_yield`

### 2.1.2 Example of Coroutines

Simple `co_await` example

2.1.3 Simple generator (`co_yield`) example

Advanced Example

### 2.1.4 Tricks and Pitfalls

## 2.2 Computer Architecture, Cache Misses

2.2.1 Different Computer Architectures, x86-64, amd, apple, mips

2.2.2 Introduction to computer architecture and storage layout

2.2.3 What are cache misses?

2.2.4 When can they happen?

2.2.5 Why are they relevant in data bases and different index structures

## 2.3 Data Structures

### 2.3.1 Hashtable

Chaining vs Open Addressing (Linear Probing)

### 2.3.2 B+ Tree

## **3 Using Coroutines in different DataStructures for hiding cache-misses**

### **3.1 Different coroutine approaches and their trade-offs**

Citation test [4].

First approach was simple: create a coroutine each time from scratch for lookup. Bad performance, re-use coroutines so the creation overhead is minimal and maybe create 5-20 coroutines and reuse them for all the lookups. First with co\_yield concept and caller has to resume the coroutine, but for caller it is complex to distinguish if

### **3.2 Coroutines in chaining Hashtable**

**3.2.1 Implementation**

**3.2.2 Benchmarking and Measurements**

**3.2.3 Results**

### **3.3 Coroutines in linear probing hashtable**

**3.3.1 Implementation**

**3.3.2 Measurements**

**3.3.3 Results**

### **3.4 Coroutines in B+ Tree**

**3.4.1 Implementation**

**3.4.2 Measurements**

**3.4.3 Results**

## **4 Integration of Coroutines in DuckDB**

### **4.1 DuckDB Introduction**

### **4.2 Understanding the implemented linear Hashtable**

Citation test [4].

### **4.3 Integration of Coroutines in DuckDB**

DuckDB integration: linear probing hashtable where if same key it is chained.

Difficulty in understanding the hashtable and implementing coroutines in huge concept of hashtable without knowing full picture. has to be done in Join phase of the hashtable (which is divided in different datachunks and each datachunk is handled by a thread where a thread can possibly handle more datachunks). In this thread coroutines can easily integrated as they are asynchronously called and ...

### **4.4 Results of timings**

## **5 Evaluation**

**5.1 Coroutine Approaches evaluation**

**5.2 Coroutine in different Data structures evaluation**

**5.3 Coroutine in DuckDB evaluation**

**5.4 Summarized evaluation of coroutines in C++ for hiding  
cache misses**

## **6 Conclusion**

## **Abbreviations**

# List of Figures

2.1	Calling a Coroutine . . . . .	3
2.2	Calling a Coroutine . . . . .	3

## **List of Tables**

# Bibliography

- [1] L. Baker. *C++ Coroutines: Understanding operator co\_await*. <https://lewissbaker.github.io/2017/11/17/understanding-operator-co-await>. Accessed: 2025-10-30. Nov. 2018.
- [2] L. Baker. *C++ Coroutines: Understanding the promise type*. <https://lewissbaker.github.io/2018/09/05/understanding-the-promise-type>. Accessed: 2025-10-30. Sept. 2018.
- [3] L. Baker. *Coroutine Theory*. <https://lewissbaker.github.io/2017/09/25/coroutine-theory>. Accessed: 2025-10-30. Sept. 2017.
- [4] L. Lamport. *LaTeX : A Documentation Preparation System User's Guide and Reference Manual*. Addison-Wesley Professional, 1994.