



SCHOOL OF COMPUTATION,
INFORMATION AND TECHNOLOGY —
INFORMATICS

TECHNISCHE UNIVERSITÄT MÜNCHEN

Master's Thesis in Informatics

**Hiding Cache-Misses with Coroutines
across different Hardware Architectures**

Daniel Gruber





SCHOOL OF COMPUTATION,
INFORMATION AND TECHNOLOGY —
INFORMATICS

TECHNISCHE UNIVERSITÄT MÜNCHEN

Master's Thesis in Informatics

**Hiding Cache-Misses with Coroutines
across different Hardware Architectures**

**Verbergen von Cache-Misses mittels
Coroutinen auf unterschiedlichen
Hardwarearchitekturen**

Author:	Daniel Gruber
Examiner:	Alexander Beischl
Supervisor:	Prof. Dr. Thomas Neumann
Submission Date:	22.12.2025



I confirm that this master's thesis is my own work and I have documented all sources and material used.

Munich, 22.12.2025

Daniel Gruber

Acknowledgments

Thanks to my supervisor Alexander Beischl for his support and guidance throughout the thesis, in particular setting up infrastructure, helping with DuckDB internals and providing feedback on the writing as well as initial guidance on benchmarking setup.

Abstract

Contents

Acknowledgments	iv
Abstract	v
1 Introduction	1
1.1 Section	1
1.2 Related Work	1
1.3 Structure of this thesis	1
2 Coroutines and Cache Misses Fundamentals	2
2.1 C++ Coroutines	2
2.1.1 Introduction to Coroutines	2
2.1.2 Execution of a Coroutine	5
2.1.3 Implementation details/concepts of a coroutine	6
2.2 Cache Misses and Software Prefetching	10
2.2.1 Definition of Cache Misses	11
2.2.2 Cause of Cache Misses in Data Structures	12
2.2.3 Techniques to counter Cache Misses	12
2.2.4 Software Prefetching	14
3 Method to hide cache misses with coroutines	18
3.1 Why using Coroutines combined with Software Prefetching?	18
3.1.1 Different Methods using Coroutines	18
3.1.2 Different coroutine approaches and their trade-offs	18
3.1.3 Scheduler for multiple coroutines	20
3.2 Benchmarking Setup	21
3.2.1 Testing Coroutines in Different DataStructures with different compilers across Different Machines	21
3.2.2 Used Benchmarking Frameworks	21
3.2.3 Benchmark configuration and recorded metrics	22

Contents

4 Using Coroutines and Cache Prefetching in different DataStructures for hiding cache-misses	23
4.1 Coroutines in chaining Hashtable	23
4.2 Coroutines in linear probing hashtable	24
4.3 Coroutines in B+ Tree	26
5 Integration of Coroutines in DuckDB	27
5.1 DuckDB Overview	27
5.2 Understanding the Linear Hashtable in DuckDB	28
5.3 Integration of Coroutines in DuckDB	28
5.4 Benchmarking	28
5.4.1 Results per machine	29
5.4.2 Results per compiler	29
5.4.3 results across machines and compiler	29
5.4.4 results per query	29
6 Evaluation	30
6.1 B+ Tree evaluation	30
6.2 Hashtable evaluation	30
6.3 Linear Probing Hashtable evaluation	30
6.4 DuckDB Evaluation	30
6.5 Coroutine Approaches evaluation	30
6.6 Coroutine in different Data structures evaluation	30
6.7 Coroutine in DuckDB evaluation	30
6.8 Summarized evaluation of coroutines in C++ for hiding cache misses	30
7 Conclusion	31
Abbreviations	32
List of Figures	33
List of Tables	34
Bibliography	35

1 Introduction

1.1 Section

Citation test [5].

How can Coroutines hide latencies for hashtable/bptree lookups? What is the performance gain generally? What is the performance gain in duckdb?

How does the performance gain differ between various architectures? (Intel x86 vs ARM vs NUMA) What about Threading and Coroutines? Hyperthreading/NUMA

1.2 Related Work

boost coro, other async concepts, other works regarding coroutine implemenmtation a lot of work on coroutines microbenchmarking on one machine with one compiler but not acrossos machines and compilers

1.3 Structure of this thesis

2 Coroutines and Cache Misses Fundamentals

2.1 C++ Coroutines

2.1.1 Introduction to Coroutines

C++ Coroutines are available in the std namespace from the C++20 standard on, and from the C++23 standard there is a generator implementation based on the C++ 20 coroutine concept.

A coroutine is a generalization of a function, being able additionally to normal functions to suspended execution and resumed it later on, as illustrated in Figure 2.1. Any function is a coroutine if its definition contains at least one of these three keywords:

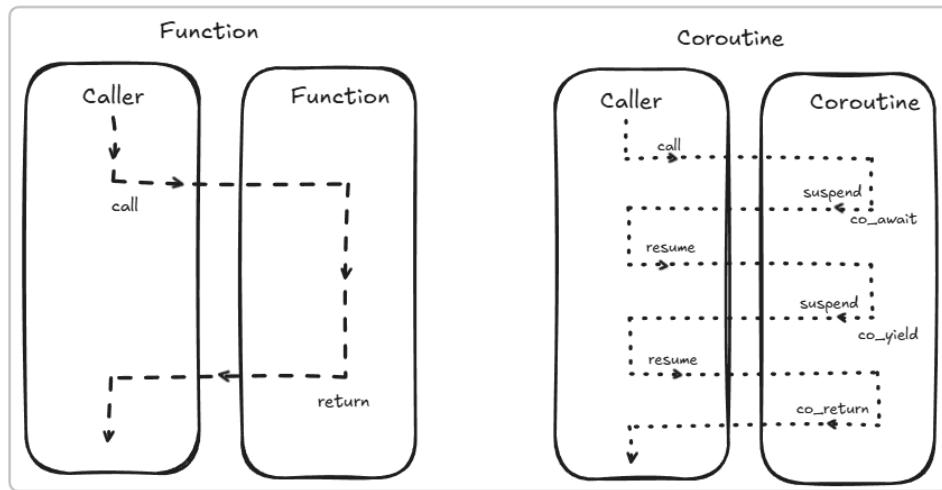


Figure 2.1: Function vs Coroutine Execution

- `co_await` - to suspend execution
- `co_yield` - to suspend execution and returning a value

- `co_return` - to complete execution and returning a value

C++ Coroutines are stackless, meaning by suspension and consecutively returning to the caller, the data of the coroutine is stored separately from the stack, namely on the heap. This allows sequential code to be executed asynchronously, without blocking the thread of execution and supports algorithms to be lazily computed, e.g. generators. However, there are some restrictions to coroutines: They cannot use variadic arguments, plain return statements or placeholder return types like `auto` or `Concept`. Also `constexpr`, `constexpr` and the `main` function as well as constructors and destructors cannot be coroutines.

To illustrate the difference between coroutines and functions, this paragraph provides a background and general information of function calls and return: A normal function has a single entry point - the Call operation - and a single exit point - the Return operation. The Call operation creates an activation frame, suspends execution of the caller and transfers execution to the callee, where the caller is the invoking function and the callee is the invoked function. The Return operation returns the value in the return statement to the caller, destroys the activation frame and then resumes execution of the caller. These operations include calling conventions splitting the responsibilities of the caller and callee regarding saving register values to their activation frames. The activation frame is also commonly called stack frame, as the functions state (parameters, local variables) are stored on the stack. Normal Functions have strictly nested lifetimes, meaning they run synchronously from start to finish, allowing the stack to be a highly efficient memory allocation data-structure for allocation and freeing frames. The pointer pointing at the top of the stack is the `**rsp**` register on X86-64 CPU Architectures.

Coroutines have, additionally to the call and return operation, three extra operations, namely suspend, resume and destroy. As coroutines can suspend execution without destroying the activation frame, as it may be resumed later, the activation frames are not strictly nested anymore. This requires that after the creation of the coroutine on the stack, the state of the coroutine is saved to the heap, like illustrated in Figure 2.2 where a normal function `f()` calls a coroutine function `c()`.

If the coroutine calls another normal function `g()`, `g()`'s activation frame is created on the stack and the coroutine stack frame points to the heap allocated frame, as illustrated in Figure 2.3. When `g()` returns, it destroys its activation frame and restores `c()`'s activation frame.

If `c()` now hits a suspension point as shown in Figure 2.5, the Suspend operation is invoked where `c()` suspends execution and returns control to `f()` without destroying its activation frame. The Suspend operation interrupts execution of the coroutine at a current, well-defined point - `co_await` or `co_yield` -, within the function and potentially transfers execution back to the caller without destroying the activation frame.

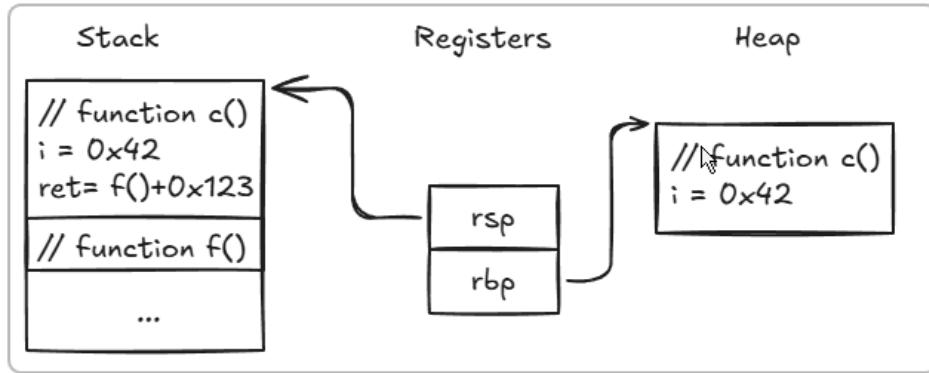


Figure 2.2: Calling a Coroutine - Heap Allocation




Figure 2.3: Coroutine calling a normal Function g()

Figure 2.4: Normal Function g() returns

This results in the stack-frame part of c() being popped off the stack while leaving the coroutine-frame on the heap. When the coroutine suspends for the first time, a return-value is returned to the caller. This return value often holds a handle to the coroutine-frame that suspended that can be used to later resume it. When c() suspends it also stores the address of the resumption-point of c() in the coroutine frame (in the illustration called RP for resume-point).

The Resume operation transfers execution back to the coroutine at the point it was suspended whereas the Destroy operation is the only operation that destroys the activation frame of the coroutine.

This handle may now be passed around as a normal value between functions. At some point later, potentially from a different call-stack or even on a different thread, something (say, h()) will decide to resume execution of that coroutine. For example, when an async I/O operation completes.

The function that resumes the coroutine calls a void resume(handle) function to

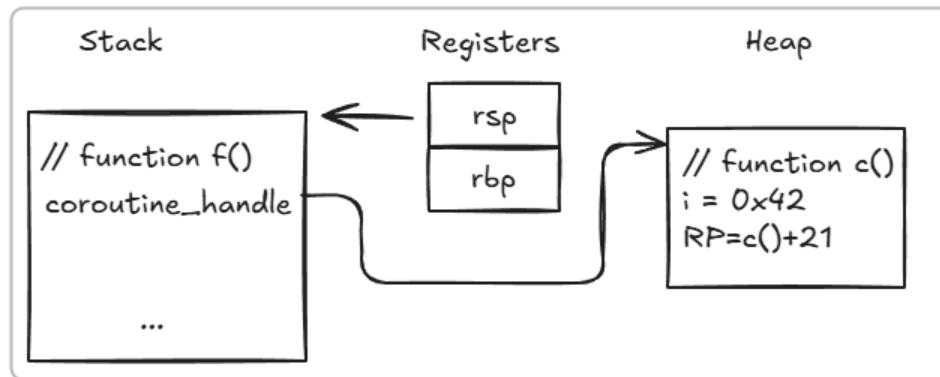


Figure 2.5: Coroutine hits Suspension Point

resume execution of the coroutine. To the caller, this looks just like any other normal call to a void-returning function with a single argument.

This creates a new stack-frame that records the return-address of the caller to resume(), activates the coroutine-frame by loading its address into a register and resumes execution of x() at the resume-point stored in the coroutine-frame.

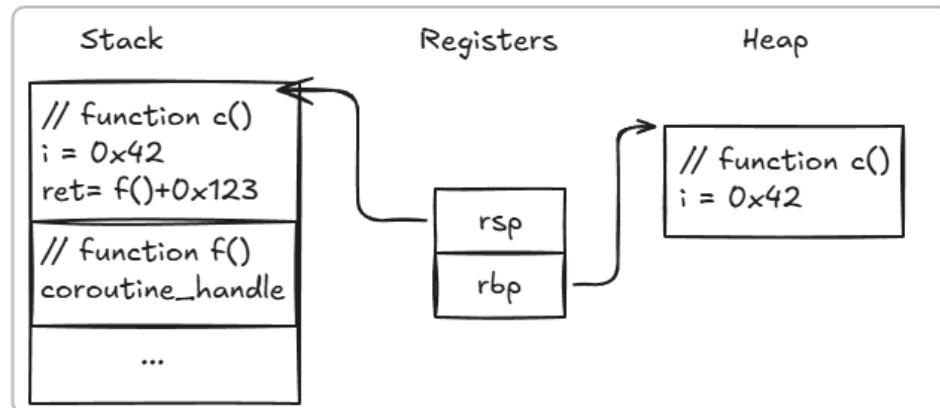


Figure 2.6: Resumption of a Coroutine

[3]

2.1.2 Execution of a Coroutine

Each coroutine is associated with a promise object, a coroutine handle and the coroutine state.

As partially illustrated in Figure 2.2, when a coroutine begins execution, it performs the following:

- allocates the coroutine state object using operator new and copies all function parameters to the coroutine state.
- calls the constructor for the promise object.
- calls promise.get_return_object() and stores the result in a local variable, which will be returned to the caller on the first suspension point of the coroutine.
- calls promise.initial_suspend() and co_await its result. Typically returns std::suspend_always (lazily-started) or std::suspend_never (eagerly-started).
- when co_await promise.initial_suspend() resumes, starts executing the body of the coroutine.

```
{
    co_await promise.initial_suspend();
    try
    {
        <body-statements>
    }
    catch (...)
    {
        promise.unhandled_exception();
    }
    FinalSuspend:
    co_await promise.final_suspend();
}
```

Listing 2.1: coroutine execution

When a coroutine reaches a suspension point, co_await or co_yield, the return object obtained earlier is returned to the caller/resumer. When encountering the co_return statement, it calls either promise.return_void() or promise.return_value(expr) depending on whether an expression is provided and whether this expression is non-void. Furthermore, all variable with automatic storage duration are destroyed in reverse order of their creation and finally calling promise.final_suspend() and co_awaiting its result. When falling off the end of the coroutine, meaning there is no co_return statement, it is equivalent to a co_return; statement.

2.1.3 Implementation details/concepts of a coroutine

The following is a minimum coroutine example using co_await, illustrated in Listing 2.2. The Task struct is a the coroutine wrapper holding the coroutine handle and the promise_type struct defining the promise type and thus, the behavior of the coroutine. The function Task Coroutine(int num_steps) is the coroutine function containing a co_await expression suspending the coroutine for num_steps times. And as explained

in previous subsection, also see Listing 2.1, when another function, e.g. main(), calls the coroutine, the coroutine frame is allocated on the heap, the promise object is constructed and initial_suspend() is called. When the coroutine hits the co_await expression, it suspends execution and returns control to the caller until it is resumed again. The coroutine results in the output of "Coroutine at step i" and "Resuming coroutine" for num_steps times, in this case from 0 to 4, as initial_suspend() is set to suspend_never and thus, the coroutine starts executing immediately when called. If it is set to suspend_always, the coroutine initially pauses before starting work (lazily computed coroutine) and results in "Resuming coroutine" being printed first followed by "Coroutine at step i" num_steps times.

```
struct Task {
    struct promise_type {
        Task get_return_object() { return Task{this}; }
        std::suspend_always initial_suspend() { return {}; }
        std::suspend_never final_suspend() noexcept { return {}; }
        void unhandled_exception() { }
        void return_void() { }
    };

    std::coroutine_handle<promise_type> h{}; // coroutine handle

    explicit Task(promise_type* p) : h{std::coroutine_handle<promise_type>::from_promised(*p)} {}
    Task(Task&& rhs) : h{std::exchange(rhs.h, nullptr)} { }
    ~Task() { if (h) { h.destroy(); } }

    static PrintCoroStep(int num_steps) {
        for (int i = 0; i < num_steps; ++i) {
            std::cout << "Coroutine\uat\ustep\u" << i << "\n";
            co_await std::suspend_always{}; // if suspend_never, coro would never pause
        }
        co_return; // equivalent to omitting this line
    }

    static main() {
        Task task = PrintCoroStep(5);
        for (int i = 0; i < 5; ++i) {
            std::cout << "Resuming\u coroutine\n";
            task.h.resume();
        }
    }
}
```

Listing 2.2: minimum coroutine co_await example

In this simple example, trivial awaitables std::suspend_always and std::suspend_never, defined in the standard library, are used to control the suspension behavior of the coroutine at the initial and co_await suspension points. The implementation of suspend_always is shown in Listing 2.3 returning false in await_ready(), indicating that the await expression always suspends and waits for a value (is not ready). The suspend_never implementation is analog with the only difference of returning true in await_ready(), indicating that the await expression never suspends (is always ready).

```
struct suspend_always {
    constexpr bool await_ready() const noexcept {return false;}
    constexpr void await_suspend(std::coroutine_handle <>) const noexcept {}
    constexpr void await_resume() const noexcept {}
};
```

Listing 2.3: suspend_always implementation

For more customization, more complex awaitable types can be implemented, requiring to implement these three methods listed in Listing 2.4.

```
bool await_ready();

// one of:
void await_suspend(std::coroutine_handle <> h);
bool await_suspend(std::coroutine_handle <> h);
std::coroutine_handle <> await_suspend(std::coroutine_handle <> h);

T await_resume();
```

Listing 2.4: awaitable type

In the *co_awaitexpr;* operator the expression is firstly converted to an awaitable type. If the promise_type type of the current coroutine has a member function await_transform, then the awaitable is obtained by calling promise.await_transform(expr), illustrated in Listing 2.5. Otherwise, the awaitable is expr as-is.

```
func get_awaitable(promise_type& promise, T&& expr){
    if P has member function await_transform:
        return promise.await_transform(expr);
    else
        return expr;
```

```

}

func get_awaiter(awaitable){
    if awaitable has member operator co_await:
        return awaitable.operator co_await();
    else if awaitable has non-member operator co_await:
        return operator co_await(awaitable);
    else:
        return awaitable;
}

```

Listing 2.5: Pseudo Code for deciding which awaier/awaitable is used

Then, the awaier object is obtained like illustrated in func get_awaiter in Listing 2.5. If no operator co_await is defined, the awaitable itself is the awaier which implicates that a type can be an awaitable and an awaier type simultaneously. For the simple example, in Listing 2.2, the expression std::suspend_always after the co_await is the expression which already is the awaitable as-is as it is an awaitable and has no await_transform() member. The awaier is also std::suspend_always as it has no operator co_await defined. Thus, std::suspend_always is both the awaitable and the awaier type.

Then, awaier.await_ready() is called , The coroutine is suspended awaier.await_suspend(handle) is called, where handle is the coroutine handle representing the current coroutine. Inside that function, the suspended coroutine state is observable via that handle, if await_suspend returns void, control is immediately returned to the caller/resumer of the current coroutine (this coroutine remains suspended), otherwise if await_suspend returns bool, the value true returns control to the caller/resumer of the current coroutine the value false resumes the current coroutine. if await_suspend returns a coroutine handle for some other coroutine, that handle is resumed (by a call to handle.resume()) Finally, awaier.await_resume() is called (whether the coroutine was suspended or not), and its result is the result of the whole co_await expr expression.

If the coroutine was suspended in the co_await expression, and is later resumed, the resume point is immediately before the call to awaier.await_resume().

Additionally to suspending the coroutine, the co_yield expression returns a value to the caller and is equivalent to co_await promise.yield_value(expr). However, to define the type to return, the promise_type has to implement the yield_value() function.

Besides the awaier/awaitable concept explained in the previous paragraphs, there is also the promise type concept. These are the two main interfaces, defined by the coroutine Type Specification (TS), for customizing the behavior of coroutines and co_await expressions:

- Awaiter / Awaitable: specifies methods that controls the semantics of `co_await` expression. When a value is `co_awaited`, the awaitable object allows to specify whether to suspend, execute some logic after suspensions (for asynchronously completed operations) and/or execute some logic after the coroutine resumes.
- Promise Type: specifies methods for customising the behavior of a coroutine, e.g. the behavior of any `co_await` or `co_yield` expression inside the coroutine body.

The promise_type struct has to be implemented in the coroutine wrapper type, which can be named arbitrarily, e.g. Task in Listing 2.2. It has to follow the scheme illustrated in Listing 2.6.

```
struct promise_type{
    // required methods
    ReturnType get_return_object();
    std::suspend_always initial_suspend();
    std::suspend_always final_suspend() noexcept;
    void unhandled_exception();

    // depending on ReturnType
    void return_void(); // if ReturnType is void
    void return_value(T value); // if ReturnType is T

    // optional methods
    auto await_transform(U&& value); // customize co_await behavior
    auto yield_value(V value); // customize co_yield behavior
}
```

Listing 2.6: Promise Type

The Promise type is determined by the compiler from the return type of the coroutine using `std::coroutine_traits`.

As mentioned in Subsection 2.1.2, each coroutine is associated with a coroutine handle. This handle is a non-owning reference to the coroutine frame and is used to resume execution of the coroutine or to destroy the coroutine frame.

[1] [2]

2.2 Cache Misses and Software Prefetching

This section provides a background on cache misses, why they cause latencies in modern computer architectures by accessing data and how they can be potentially be

hidden.

2.2.1 Definition of Cache Misses

To explain cache misses, first the underlying concept of caches in modern computer architectures has to be explained briefly. Caches are small, fast memory units located close to the CPU cores, designed to store frequently accessed data and instructions to reduce the average time to access memory. When the CPU needs to read or write data, it first checks if the data is present in the cache (a cache hit). Typically, there are three levels of caches (L1, L2, L3 or called LLC) with L1 being the smallest and fastest, and L3 being the largest and slowest. If the data is not found in the caches (a cache miss), it has to be fetched from the slower main memory.

As described, a cache miss is a failed attempt to read or write a piece of data in the cache, which results in a main memory access with much longer latency. Data is transferred between memory and cache in fixed-size blocks called cache lines. When a cache line is requested, the cache checks if it contains the data at the requested memory location (a cache hit). If not (a cache miss), the cache allocates a new entry and fetches the data from main memory. Cache misses cause stalls because the CPU must wait for data to arrive from the slower main memory. Modern CPUs can execute hundreds of instructions in the time it takes to fetch a single cache line, making these stalls particularly costly. To mitigate this, modern architectures employ techniques such as out-of-order execution, which allows the CPU to execute independent instructions while waiting for cache miss data, and simultaneous multithreading (SMT), which enables an alternate thread to utilize the CPU core during the stall. There are three types of cache misses: instruction read miss, data read miss, and data write miss. Instruction cache misses typically incur the largest penalty, as the processor must stall until the instruction is fetched from main memory. Data cache read misses cause smaller delays because independent instructions can continue execution while waiting for data and Data cache write misses result even in the shortest delays, as writes can be queued without blocking subsequent instructions.

In this work, the focus lies on data read misses, as they are the most common type of cache misses encountered during data structure traversals. These data read cache misses are typically categorized into three main types: compulsory misses, capacity misses, and conflict misses. Compulsory misses occur on the first access to a block, as the data has not yet been loaded into the cache regardless of cache design. Capacity misses happen when the cache can not hold all the data actively used by a program during execution, meaning the working set exceeds the cache size. Conflict misses, also known as interference misses, occur when multiple memory addresses map to the same cache set, leading to evictions of useful data even though the cache is not

full. For a multi-processor system this 3Cs group of cache misses can be extended to 4Cs. The 4th C are coherence misses arising from a cache line being invalidated due to modifications made by other processors (for L3 Cache as it is typically shared across cores) or threads (for L1, L2 and L3 Cache).

In this work, the focus lies merely on the first two Cs, compulsory and capacity misses, as the first are inherent to the initial data access and the second can be mitigated through better data locality and prefetching techniques.

2.2.2 Cause of Cache Misses in Data Structures

Modern database engines utilize various in-memory data structures that are directly addressed by virtual memory pointers. Memory blocks used by these data structures are typically allocated from the heap and chained together using pointers. For example, the nodes of an in-memory B+-tree are dynamically allocated and deallocated as the tree grows and shrinks. To traverse such a tree from its root node to the target leaf node, the accessing thread must dereference multiple pointers sequentially from the root to the leaf, forming a random access pattern. Another example is a chaining Hashtable, where each bucket contains a linked list of entries that hash to the same index. To look up a key in such a hashtable, the accessing thread first computes the hash index, accesses the corresponding bucket, and then traverses the linked list by following pointers until it finds the desired key or reaches the end of the list.

This pointer chasing behavior poses a significant challenge for hardware prefetchers. Such access patterns are very difficult, if not impossible, for hardware prefetchers to predict accurately, leading to a high likelihood of last-level cache misses upon each pointer dereference. But this pointer chasing is also for software prefetching challenging, as the address of the next memory access is not known until the current access is completed and the node or linked list entry is fetched and processed to extract the next pointer.

2.2.3 Techniques to counter Cache Misses

The literature presents various software techniques to address cache misses. Based on their impact on the number of cache misses and the incurred penalty, these techniques can be categorized into three main approaches.

- Eliminate cache misses by increasing spatial and temporal locality.
- Reduce the cache miss penalty by scheduling independent instructions to execute after a load.
- Hide the cache miss penalty by overlapping memory accesses.

The first approach improves cache locality by either eliminating indirection by designing cache-conscious data structures, matching the data layout to the access pattern of the algorithm such that data accessed together is stored in contiguous space, or reorganizing memory accesses. However, these optimizations often require fundamental changes to data structure designs.

The second approach focuses on reducing the cache miss penalty by scheduling independent instructions to execute after a memory load. This increases instruction-level parallelism and leads to more effective out-of-order execution. Finally, to reduce the main-memory access penalty, a non-blocking load, a prefetch instruction, must be introduced early enough to allow independent instructions to execute while fetching data. However, this approach relies on the presence of independent instructions following the load and timing it correctly to ensure that the data is available when needed.

The third approach seeks to hide the cache miss penalty by overlapping memory accesses, meaning issuing multiple memory requests in parallel. This, however, requires independent memory accesses which are not naturally available in dependent access patterns like index lookups.

These traditional approaches provide limited benefits for individual index lookups with dependent memory accesses if parallelism cannot be exploited. To effectively hide cache misses in such scenarios, techniques that can create parallelism across multiple independent operations are needed, such as processing operations in groups with interleaved execution.

In the literature two main techniques have been proposed to achieve this, :

- Group Prefetching (GP)
- Asynchronous Memory Access Chaining (AMAC)
- Coroutines

Group Prefetching (GP) interleaves N identical operations across M predefined code stages, executing each operation through all stages sequentially before moving to the next stage. While GP maintains a "semi-synchronous" programming model, it has significant limitations: it requires knowing the number of stages in advance, and cannot efficiently handle early-terminating operations, leading to performance degradation on irregular access patterns.

Asynchronous Memory Access Chaining (AMAC) transforms operations into state machines stored in a circular buffer, allowing different operations to execute different stages concurrently. This enables immediate reuse of completed operations without waiting for others. However, AMAC requires complete code rewriting, resulting in code that is difficult to understand and maintain.

Both GP and AMAC achieve high performance through careful hand-coding, but at the cost of developer productivity. Coroutines offer an alternative that combines the benefits of both approaches: they enable interleaved execution like AMAC while preserving the synchronous programming model like GP, but with compiler-generated context-switching code rather than manual state machine implementation. Developers simply add suspension points (`co_await` or `co_yield`) after prefetch operations, and the compiler handles the efficient state management and context switching.

All these techniques rely on using software prefetching to initiate memory transfers ahead of time which is explained in the next subsection.

2.2.4 Software Prefetching

Hardware manufacturers have implemented various hardware prefetching techniques to automatically detect and pre-load data into caches before they are accessed by the CPU. These hardware prefetchers monitor memory access patterns and attempt to predict future accesses based on historical data. However, the reliance on hardware prefetchers has limitations, especially for applications with irregular access patterns where predictions may be inaccurate or insufficient.

The solution to overcome this is Software-based prefetching. It enables applications to proactively hint to the CPU about upcoming data accesses which for the hardware is hard or even impossible to predict. These hints, specific instructions in modern ISAs such as x86 and ARMv8, allow the hardware to move data asynchronously into the cache hierarchy, thereby reducing access latency by overlapping data transfers with computation.

With Software prefetching, proactively hinting the CPU what data to preload into the cache, timing of these prefetches is crucial for effectiveness. When software executes the prefetching instruction too early before the actual access, the data may have already been evicted from the cache when needed. Conversely, if prefetched too late, the data might not be transferred to the cache completely, potentially causing the CPU to stall. The presence of intricate memory systems like NUMA and high-bandwidth memory, as well as the increasing sophistication of CPUs with out-of-order execution and superscalar architectures, further complicate achieving precise timing.

Before diving into the details of software prefetching, it is essential to understand the data load process through the memory subsystem of modern computer architectures. As illustrated in Figure 2.7, when the CPU executes a load instruction to read data from memory, it first translates the given virtual address into a physical address using the Memory Management Unit (MMU), in particular the Translation Lookaside Buffer (TLB). If the TLB does not contain the mapping for the virtual address, a page-table walk is initiated to retrieve the mapping from the page tables in memory and to finally

update the TLB. Once the physical address is determined, the CPU checks the cache hierarchy (L1, L2, L3) for the requested data. If the data is found in any of the cache levels (a cache hit), it is returned to the CPU. However, if the data is not present in the caches (a cache miss), the CPU issues a request to the main memory to fetch the data.

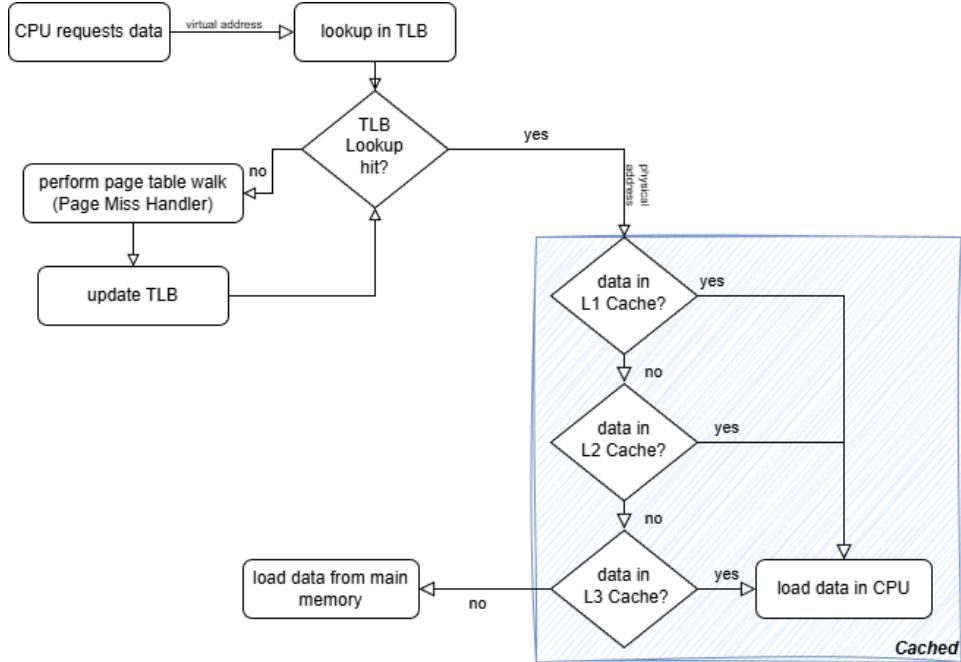


Figure 2.7: Data Load through Memory Subsystem

Back to prefetching, modern platforms offer several prefetch instructions to target different levels of the cache hierarchy, with each instruction initiating the transfer of one cache line. As visualized in Figure 2.8, when the software executes a prefetch, the logical address is first translated into a physical one. This procedure happens synchronously—if the translation is not cached within the Translation Lookaside Buffer (TLB), the CPU contacts a Page Miss Handler (PMH) before continuing execution of the prefetch instruction.

Once the physical address is known and the request misses the L1 data cache (L1d), the hardware triggers the transfer from memory to caches asynchronously. The L1d cache requests the cache line through the Line Fill Buffer (LFB), which acts as a buffer-like structure interfacing between the L1d and L2 cache to communicate requests on a cache-line granularity. For every cache line that misses the L1d, the CPU allocates a slot in the LFB. This nature makes it straightforward to implement asynchronous prefetches: the desired cache line does not cause the instruction to wait until the data is transferred

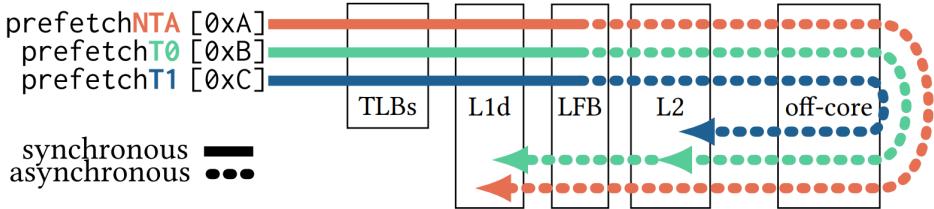


Figure 2.8: Lifecycle of prefetch instructions through the memory subsystem

into the cache but only until the miss is communicated to the LFB. Furthermore, the LFB enables the CPU to handle several pending requests simultaneously without blocking, supporting out-of-order execution and performing micro-optimizations like merging multiple requests to the same cache line.

The specific software prefetch instruction dictates how close the data is moved to the CPU. In the x86 ISA, most prefetch instructions clearly specify the cache level target for the fetched data: `prefetcht1` stores cache lines in the L2 cache, `prefetcht0` moves data into the L1d, and `prefetcht2` is generally aimed at the LLC (though on recent Intel platforms since Skylake with non-inclusive victim caches, data is placed into the L2 instead). The non-temporal `prefetchnta` instruction does not specify a particular cache target but aims to reduce cache pollution for data that will be accessed non-recurrently, with these cache lines being prioritized for quicker eviction.

Modern platforms exhibit two notable limitations in the implementation of software prefetches. First, the execution of a single prefetch instruction stalls until the virtual address is translated into a physical address. Although the TLB might accelerate this process, prefetched addresses are often not present in the TLB since software prefetching is primarily used when the application accesses scattered data objects. Consequently, the latency of the prefetch instruction is dominated by the latency of the PMH performing a page-table walk to translate the address. However, this restriction only applies to the initial prefetch operation within a memory page if several prefetches are performed consecutively.

Second, due to the limited number of LFB entries (typically between 3 and 24), the hardware can accommodate only a small number of outstanding memory requests. If an excessive volume of prefetch requests is issued rapidly, the LFB becomes saturated and cannot accept further requests. Under this circumstance, the CPU will either stall until an LFB slot becomes available or drop new prefetch requests. This limitation necessitates careful management of prefetch intensity to balance memory-level parallelism with hardware constraints.

Thus, to efficiently exploit software prefetching in data structures like trees and hash

tables, these prefetch instructions must be strategically placed within the code to have the data in cache when needed. As mentioned in previous chapter, techniques such as breaking operations into stages and pipelining, or utilizing asynchronous control flow abstractions like coroutines, can facilitate effective prefetching by creating sufficient temporal gaps between prefetch issuance and data access.

3 Method to hide cache misses with coroutines

This chapter describes the history of different methods that were used to implement coroutines for hiding cache misses in data structures. It also explains why coroutines and cache prefetching were chosen as the method to hide cache misses in pointer chasing data structures and which method is superior to others in terms of performance and implementation complexity.

3.1 Why using Coroutines combined with Software Prefetching?

This thesis focuses on using coroutines and cache prefetching to hide cache misses in pointer chasing data structures. The reason for this choice of using coroutines (see Section 2.1) and software prefetching (Subsection 2.2.4) is, that coroutines offer an alternative that combines the benefits of AMAC and GP approaches: they enable interleaved execution like AMAC while preserving the synchronous programming model like GP, but with compiler-generated context-switching code rather than manual state machine implementation. Suspension points like `co_await` or `co_yield` can be simply added after prefetch operations, and the compiler handles the efficient state management and context switching by transforming the coroutine function into a state machine.

As Coroutines are already widely used in asynchronous programming to handle I/O-bound tasks without blocking threads successfully in C++ applications, this thesis explores their potential in hiding cache misses in pointer-chasing data structures across machines and compilers.

3.1.1 Different Methods using Coroutines

3.1.2 Different coroutine approaches and their trade-offs

During the development of coroutine-based data structure lookups, the method of implementing coroutines evolved through several iterations. The following paragraphs

describe the different approaches taken, their trade-offs, and the reasoning behind the final chosen method.

The first approach used is the most straightforward one, which is creating a new coroutine instance for each lookup value/operation. This means, each time a lookup is performed, a new coroutine is instantiated, executed until the first suspension point (after prefetching), and then destroyed once the lookup is complete. This involves significant overhead due to the frequent creation and destruction of coroutine state machines, which includes allocating stack space and managing the coroutine's context, including saving to the heap as explained in Section 2.1. Despite being the method most easiest to understand, having to do all this overhead implicates that this method is inferior in performance compared to other methods.

This major drawback leads to the second approach, which is to create a fixed number of coroutines at the beginning of the lookup process and to re-use coroutines so the creation overhead is minimal and the lookup process can benefit from interleaved execution of multiple lookups. With this second approach, the implemented coroutine function for the lookup could now leverage `co_yield` to return found values back to the caller, while the caller is responsible for resuming the coroutine until all lookups are processed. Furthermore, the caller is responsible for distinguishing between three states:

- 1) the lookup is not yet finished (the coroutine yielded without finding the value yet),
- 2) the lookup found the value (the coroutine yielded with a found value), and
- 3) the lookup did not find the value (the coroutine finished without yielding a value).

Looking at this, it is not a binary choice between found vs. not found, but a ternary choice that the caller has to handle. To distinguish between these three states, as solution, the coroutine function could yield an optional value type, where a value indicates that the lookup found the value, and a `nullopt` indicates that the lookup is not yet finished. However, this requires either a fundamental change to the data structures implementation, having a `nullopt` value type, or constructing a wrapper type that can hold either a valid value or a `nullopt` making the lookup procedure more complex.

This complexity and the need to handle three different states led to the final approach, which is to pass a pointer/reference to the result value into the coroutine function. The coroutine function then writes the found value into the provided pointer/reference when the lookup is successful. If the lookup is not yet finished, the coroutine simply suspends without modifying the pointer/reference.

3.1.3 Scheduler for multiple coroutines

The presented scheduler in [4], illustrated in Listing 3.1

```
for i = 0 to batch_size - 1:  
    coroutine[i] = foo(...);  
while any(coroutine_promises, x: not x.done()):  
    for i = 0 to batch_size - 1:  
        if not coroutine[i].done():  
            coroutine[i].resume()
```

Listing 3.1: CoroBase Scheduler for lookup

```
template<const size_t numCoroutines>  
std::vector<typename resultType>  
lookupEntries(datastructure ds, lookupsVector lookups) noexcept {  
    resultType result;  
    std::array<Silent, numCoroutines> coroutines =  
        ([&ht, &lookups, &result]<size_t... Is>(std::index_sequence<Is...>) {  
            return std::array{lookup<numCoroutines>(ht, lookups, result, Is)...};  
        })(std::make_index_sequence<numCoroutines>{});  
    while (true) {  
        size_t activeCoroutines = 0;  
        for (auto &c: coroutines) {  
            if (!c.h_.done()) {  
                activeCoroutines++;  
                c.h_.resume();  
            }  
        }  
        if (activeCoroutines == 0) { break; }  
    }  
    return result;  
}
```

Listing 3.2: Scheduler for lookup

3.2 Benchmarking Setup

3.2.1 Testing Coroutines in Different DataStructures with different compilers across Different Machines

3.2.2 Used Benchmarking Frameworks

To evaluate the coroutine implementations, two benchmarking setups are used: Google Benchmark and the perf-cpp framework developed at TU Dortmund. Google Benchmark offers a quick way to prototype individual microbenchmarks and is therefore used early on to debug coroutine state machines, validate correctness of lookups under different prefetch distances, and sanity-check throughput improvements on a small scale. Its minimal configuration effort and familiar API is convenient for short, focused experiments, especially in the beginning when iterating over scheduler changes or changes in the coroutine lookup functions themselves. Despite its convenience, Google Benchmark has limitations in terms of flexibility for constructing complex benchmark scenarios that involve multiple data structures, shared setup code, and detailed performance measurements using hardware counters. Even though, fixtures are available in Google Benchmark, they lack the capability to share complex setup logic across multiple benchmarks (e.g., building a hash table once and reusing it across different lookup benchmarks with varying coroutine parameters) and using templated functions evaluated at build time to generate multiple benchmark instances with different configurations.

For the final evaluation, perf-cpp is used as the primary framework because it provides richer control over benchmark composition, data structure construction and re-use across different benchmarks, and parameter sweeps. The framework allows building reusable benchmark fixtures that set up hash tables or trees once and then run multiple lookup benchmarks against the same data structures. Additionally, perf-cpp supports detailed performance measurements using hardware counters, enabling fine-grained analysis of cache misses, branch mispredictions, and other low-level metrics relevant to evaluating coroutine performance.

This flexibility was essential to compare different coroutine designs (e.g., varying the number of reused coroutine instances or altering prefetch strategies) under consistent data and workload characteristics without duplicating setup code.

In conclusion, Google Benchmark serves as a pointwise verifier, where all the benchmark scenarios are mirrored but with less flexibility in terms of choosing datasize, reusing datastructures for several different lookup functions and setting several parameters at once. However, all reported measurements and conclusions in this thesis rely on perf-cpp runs, as that framework proved more expressive and better suited for constructing realistic, repeatable experiments across the evaluated data structures.

3.2.3 Benchmark configuration and recorded metrics

Different number of corouitnes, using with threading, using with different data sizes, prefetch distances, ...

4 Using Coroutines and Cache Prefetching in different DataStructures for hiding cache-misses

This chapter focuses on the implementation and evaluation of coroutines in various pointer-chasing data structures, including chaining hash tables, linear probing hash tables, and B+ trees. It explains how coroutines are integrated into these data structures to hide cache misses during lookups and presents benchmarking results to assess their performance impact.

4.1 Coroutines in chaining Hashtable

```
template<const size_t numCoroutines>
Task lookup(const ht_primary_key::Hashtable &ht, const auto &lookups, auto &results,
           const size_t offset) noexcept {
    auto idx = offset; size_t key; uint64_t hashIdx;

    for (; idx < lookups.size(); idx += numCoroutines) {
        key = lookups[idx];
        hashIdx = hashKey(key) & ht.mask;
        PREFETCH(&(ht.ht[hashIdx])); // Initial prefetch of the first value
        co_await std::suspend_always{}; // Suspend the coroutine until resumed

        auto result = ht.ht[hashIdx]; // Read the previously cached value

        while (!((result == nullptr) || (result->key == key))) {
            PREFETCH(result->next); // Prefetch the next entry in the linked list
            co_await std::suspend_always{};
            result = result->next;
        }
        if (result != nullptr) results.push_back(result);
    }
}
```

```
    co_return;
}
```

Listing 4.1: coroutine lookup in chaining hashtable

4.2 Coroutines in linear probing hashtable

```
template<const size_t numCoroutines>
Task lookup_ht_entry_t(const linear_probing_ht_duckdb::LinearProbingHashTable<uint64_t>
                      const auto &lookups, auto &results, const size_t offset) noexcept {
    auto idx = offset;
    size_t key;
    uint64_t hashIdx;
    uint64_t hash;
    linear_probing_ht_duckdb::LinearProbingHashTable<uint64_t, uint64_t>::EntryData *data;

    for (; idx < lookups.size(); idx += numCoroutines) {
        key = lookups[idx];
        hash = ht.ComputeHash(key);
        uint64_t slot = hash & ht.bitmask;
        uint64_t salt = ht.ExtractSalt(hash);

        // Find SLOT
        linear_probing_ht_duckdb::ht_entry_t entry;
        while (true) {
            PREFETCH(&(ht.entries[slot])); // Initial prefetch of the first value
            co_await std::suspend_always{}; // Suspend the coroutine until resumed
            entry = ht.entries[slot]; // Read the previously cached value
            if (!entry.IsOccupied()) {
                // Empty slot found
                break;
            }
            if (ht.use_salts) {
                // Check salt first (bloom filter)
                if (entry.GetSalt() == salt) {
                    // Salt matches, need to check actual key
                    data = reinterpret_cast<
                        linear_probing_ht_duckdb::LinearProbingHashTable<uint64_t, ui
```

```

        entry.GetPointer());
    if (data && ht.key_equal(data->key, key)) {
        entry = ht.entries[slot]; // Found matching key
        break;
    }
}
} else {
    // No salt optimization, check key directly
    data = reinterpret_cast<
        linear_probing_ht_duckdb::LinearProbingHashTable<uint64_t, uint64_t>::EntryData*
    >(entry.GetPointer());
    if (data && ht.key_equal(data->key, key)) {
        entry = ht.entries[slot]; // Found matching key
        break;
    }
}

// Continue linear probing
data = nullptr;
linear_probing_ht_duckdb::IncrementAndWrap(slot, ht.bitmask);
}

if(data == nullptr) {
    continue; // Not found
}
// Found an entry, retrieve all values in the chain
linear_probing_ht_duckdb::LinearProbingHashTable<uint64_t, uint64_t>::EntryData* current;
while (current) {
    results.push_back(current);
    PREFETCH(current->next);
    co_await std::suspend_always();
    current = current->next;
}
}

co_return;
}

```

Listing 4.2: linear probing hashtable coroutine lookup

4.3 Coroutines in B+ Tree

```

template<typename Key_t, typename Value_t, const size_t numCoroutines>
Task lookup(const bptree::BPTree<Key_t, Value_t> &t, const auto &lookups, auto &result
           const size_t offset) noexcept {
    auto idx = offset;
    uint64_t key;

    for (; idx < lookups.size(); idx += numCoroutines) {
        key = lookups[idx];

        auto node = t.root;
        while (!node->isLeaf()) {
            size_t i = 0;
            while (i < node->nr_keys && key >= node->keys[i]) {
                i++;
            }
            PREFETCH(std::move(dynamic_cast<bptree::PtrNode<Key_t, Value_t> *>(node)->children[i]));
            co_await std::suspend_always{};
            node = dynamic_cast<bptree::PtrNode<Key_t, Value_t> *>(node)->children[i];
        }
        bptree::LeafNode<Key_t, Value_t> *leafNode =
            std::move(dynamic_cast<bptree::LeafNode<Key_t, Value_t> *>(node));

        // here I could do better search algorithm (binary search) but for now I just do linear search
        for (size_t i = 0; i < leafNode->nr_keys; i++) {
            if (leafNode->keys[i] == key) {
                results.push_back(std::make_tuple(std::move(key), std::move(leafNode->values[i])));
                break; // Stop searching if key is found
            }
        }
    }

    co_return;
}

```

Listing 4.3: B+ Tree lookup coroutine

5 Integration of Coroutines in DuckDB

5.1 DuckDB Overview

DuckDB is a relational embeddable analytical DBMS that focuses on supporting analytical query workloads (OLAP). Similar to SQLite, DuckDB prioritizes simplicity and ease of integration by eliminating external dependencies for compilation and run-time. DuckDB's columnar-vectorized query execution engine reduces the CPU cycles expended per individual value, processing large batches of values in one operation as a vector. This design choice optimizes DuckDB for analytical queries, and in embedded scenarios, it enables high-speed data transfer to and from the database. To provide transactional guarantees (ACID properties), DuckDB employs its custom bulk-optimized multi-version concurrency control (MVCC) and enables data to be stored in persistent, single-file databases. Additionally, DuckDB supports secondary indexes to speed up queries trying to find a single table entry.

JOINS DuckDB supports several join algorithms, including hash join, sort-merge join, and index join. The system also implements join ordering optimization using dynamic programming and a greedy fallback for complex join graphs. It performs flattening of arbitrary subqueries and has a set of rewrite rules to simplify the expression tree, including common subexpression elimination and constant folding. The result is an optimized logical plan for the query, which is then transformed by the physical planner into the physical plan, selecting suitable implementations where applicable. In addition, DuckDB introduced a join algorithm called range join, which uses the interval encoding join (IEJoin) technique to handle range queries efficiently. IEJoin leverages min-max indexes to reduce the amount of data that needs to be scanned and uses SIMD instructions to further improve performance. Finally, DuckDB also supports an out-of-core hash-join for large tables that cannot fit in memory.

QUERY COMPILATION DuckDB does not support Just-in-Time (JIT) compilation of SQL queries, but instead chooses to use a vectorized model. This is because JIT engines require large compiler libraries such as LLVM, which can create additional transitive dependencies and make the deployment and integration process more complex. DuckDB aims to simplify this process by avoiding external dependencies and using templates for code generation.

QUERY EXECUTION DuckDB uses a vectorized interpreted execution engine that

supports fixed-length types like integers as native arrays and variable-length values like strings as a native array of pointers into a separate string heap. To represent NULL values, DuckDB uses a separate bit vector that is only present if NULL values appear in the vector. DuckDB supports a vectorized pull-based model that is also known as "vector volcano". In this model, query execution starts by pulling the first chunk of data from the root node of the physical plan. A chunk is a horizontal subset of a result set, query intermediate, or base table. This node will recursively pull chunks from child nodes, eventually arriving at a scan operator that produces chunks by reading from the persistent tables. The process continues until the chunk arriving at the root is empty, indicating that the query has been completed. DuckDB uses a vectorized processing model that processes batches of columns at a time, which is optimized for CPU cache locality, suitable for SIMD instructions and pipelining, and has small intermediates that ideally fit in L1 cache. The system contains an extensive library of vector operations that support the relational operators, and this library expands code for all supported data types using C++ code templates.

5.2 Understanding the Linear Hashtable in DuckDB

5.3 Integration of Coroutines in DuckDB

DuckDB integration: linear probing hashtable where if same key it is chained.

Difficulty in understanding the hashtable and implementing coroutines in huge concept of hashtable without knowing full picture. has to be done in Join phase of the hashtable (which is divided in different datachunks and each datachunk is handled by a thread where a thread can possibly handle more datachunks). In this thread coroutines can easily be integrated as they are asynchronously called and ...

5.4 Benchmarking

which benchmarking used, integrated benchmarking framework of duckdb and only tpch queries

5.4.1 Results per machine

5.4.2 Results per compiler

5.4.3 results across machines and compiler

5.4.4 results per query

6 Evaluation

6.1 B+ Tree evaluation

6.2 Hashtable evaluation

6.3 Linear Probing Hashtable evaluation

6.4 DuckDB Evaluation

6.5 Coroutine Approaches evaluation

6.6 Coroutine in different Data structures evaluation

6.7 Coroutine in DuckDB evaluation

**6.8 Summarized evaluation of coroutines in C++ for hiding
cache misses**

7 Conclusion

Abbreviations

List of Figures

2.1	Function vs Coroutine Execution	2
2.2	Calling a Coroutine - Heap Allocation	4
2.3	Coroutine calling a normal Function g()	4
2.4	Normal Function g() returns	4
2.5	Coroutine hits Suspension Point	5
2.6	Resumption of a Coroutine	5
2.7	Data Load through Memory Subsystem	15
2.8	Lifecycle of prefetch instructions through the memory subsystem	16

List of Tables

Bibliography

- [1] L. Baker. *C++ Coroutines: Understanding operator co_await*. <https://lewissbaker.github.io/2017/11/17/understanding-operator-co-await>. Accessed: 2025-10-30. Nov. 2018.
- [2] L. Baker. *C++ Coroutines: Understanding the promise type*. <https://lewissbaker.github.io/2018/09/05/understanding-the-promise-type>. Accessed: 2025-10-30. Sept. 2018.
- [3] L. Baker. *Coroutine Theory*. <https://lewissbaker.github.io/2017/09/25/coroutine-theory>. Accessed: 2025-10-30. Sept. 2017.
- [4] Y. He, J. Lu, and T. Wang. “CoroBase: Coroutine-Oriented Main-Memory Database Engine.” In: *Proc. VLDB Endow.* 14.3 (2021), pp. 431–444. DOI: 10.14778/3430915.3430932.
- [5] L. Lamport. *LaTeX : A Documentation Preparation System User's Guide and Reference Manual*. Addison-Wesley Professional, 1994.